

CHANDRAKANTH PERAVELLI

Chaska, Minnesota

(952) 255-9688 | chandrakanthperavelli@gmail.com | linkedin.com/in/chandrakanthperavelli

Lead / Staff Java Full Stack Engineer

PROFESSIONAL SUMMARY

Versatile Staff/Lead Software Engineer with 14+ years' experience, specializing in a design-first approach. Proficient in architecting and building data-intensive, highly available distributed systems using the Java technology stack, open-source data platforms, and multi-cloud environments.

- Extensive experience in designing and developing software applications using Java, Spring Boot, Angular, React, Node.js, RESTful APIs, Python, FastAPI, Fastify, Redux, Kotlin, Kotlin Coroutines, Kafka, MongoDB, Postgres, Elasticsearch, Redis, Docker, Kubernetes, and cloud platforms such as AWS and GCP.
- Own architecture and reliability strategy for a widely-deployed, large-scale distributed data-governance platform, driving initiatives in high availability, cloud replication, and runtime scalability that materially improved uptime and performance.
- Knowledge of data modeling, writing complex SQL queries, stored procedures, functions, and SQL optimization within an Oracle and/or SQL Server environment.
- Containerize microservices, build Docker images, push to Docker Hub, and deploy to GCP.
- Deploying multi-container microservices applications with automated CI/CD pipelines (GitHub Actions, Vela, Jenkins, Spinnaker).
- Proficient in creating and securing APIs via JWT tokens and OAuth.
- Integrate AWS services including S3, Kinesis, SNS, and SQS in serverless projects.
- Strong knowledge of OOP principles; implemented dependency injection, inversion of control, repository, unit of work, singleton, CQRS, Facade, and other proven design patterns across multiple projects.
- Extensive experience with SaaS platform infrastructure and architecture, with a focus on high availability and scalability.
- Good knowledge of MockK, JUnit automated testing, and stress/performance testing with frameworks like Spock, K6, and Cucumber.
- Coach and mentor technical resources, providing insight into architectural decisions and trade-offs.
- Lead application development/engineering teams across local and remote resources, partnering with Service Delivery Managers and Program/Project Managers.
- Communicate ideas, issues, status, and outcomes with the team, product owners, and other stakeholders.

CORE COMPETENCIES

Java | J2EE | Kotlin | Kotlin Coroutines | Groovy | Python | Go-Lang | Apache Atlas | Metadata Management | Distributed Systems | High Availability | Grafana | Kibana | Splunk | Jenkins | Cloud Computing (GCP/AWS) | React JS | Git | Docker | Kubernetes | Spring Boot | Micronaut | FastAPI | Fastify | TypeScript | Node.js | ELK Stack | Agile Methodologies | Kafka | Redis | MongoDB | DynamoDB | REST API | Postgres | jOOQ | Cucumber | K6 | ZooKeeper | JIRA

EDUCATION

Anna University, India — Bachelor of Engineering in Information Technology, May 2011 (GPA: 8/10)

PROFESSIONAL EXPERIENCE

Cloudera, Chaska, Minnesota (Remote)

[Month Year] – Present

Staff Software Engineer

- Own architecture and technical direction for a core enterprise data-governance platform relied on by large-scale customer deployments, driving reliability and scalability initiatives from design through production rollout.
- Led a major platform reliability initiative that re-architected how the system handles failover and continuous availability, eliminating a legacy infrastructure dependency and setting the platform up for near-zero downtime operation — work now treated as the blueprint for future resiliency efforts.
- Architected a next-generation modular runtime framework (AMRA) to scale the platform's distributed processing capacity, unlocking significantly higher data throughput and positioning the system to support enterprise-scale workloads well beyond its prior limits.

- Redesigned the platform's cloud-to-cloud data replication architecture, converting a synchronous process into an asynchronous, fault-tolerant pipeline — substantially improving replication performance while preserving strict data integrity guarantees at scale.
- Acted as the primary technical authority translating ambiguous, org-level reliability goals into a concrete execution plan, producing the design specification and a fully scoped delivery roadmap (21 stories across 9 epics) that the broader team executed against.
- Spearheaded a multi-quarter runtime modernization program spanning several release trains, aligning engineering, QE, and program management across teams to protect a critical release commitment through a mid-cycle scope change — delivered with zero schedule slippage.
- Recognized across engineering leadership as a force multiplier — mentoring engineers, setting architectural standards, and driving distributed-systems performance improvements across previously siloed teams.

Target Corporation, Austin, Texas

April 2018 – [Month Year]

Lead Engineer

- Architected and led the development of an enterprise-grade 3D automation pipeline, utilizing cutting-edge technologies including Airflow, Node.js, and Python, resulting in savings of approximately \$2M between Target and its vendors.
- Designed and orchestrated multiple microservices architectures, seamlessly integrating Kafka for real-time data streaming with AWS S3, GCP, Spring Boot, Kotlin, Postgres, and MongoDB backing the 3D automation system.
- Crafted algorithmic solutions to enhance guest online order fulfillment using PySpark, Python, Spring Boot, Kafka, microservices, and an event-driven framework, resulting in annual savings of over \$1M.
- Engaged with major carriers in the US to simulate performance impacts akin to holiday periods and proposed changes to handle over 1,000+ requests per second.
- Developed an enterprise-level, language-agnostic security package to manage security and token validation with OAuth, leveraging Spring Boot, Flask, and Node.js frameworks.
- Pioneered an enterprise-level Target search engine tailored for retrieving Target products based on metadata, leveraging Elasticsearch's fuzzy, exact, and autocomplete capabilities.
- Implemented automated functional and performance testing within the CI/CD pipeline, streamlining developer workflows and facilitating maintenance and creation of high-performing code.

Best Buy, Minneapolis, Minnesota

May 2015 – April 2018

Senior Software Engineer

- Developed the Best Buy Enterprise Microservices Library to mitigate API redundancy across numerous applications managing identical datasets, resulting in a 30% reduction in developer effort.
- Developed custom dashboards and metrics in Grafana/Kibana to view error reports and access logs.
- Designed Swagger UI for API testing and documentation using Swashbuckle and OpenAPI principles.
- Developed intuitive UI screens and integrated state management with React Redux to enhance user experience and streamline data flow in web applications.
- Optimized database performance using PostgreSQL and implemented efficient data access with jOOQ, ensuring rapid query execution and enhanced system responsiveness.
- Implemented and managed cloud-based solutions on AWS, leveraging various services to meet project requirements.
- Conducted thorough and insightful code reviews, fostering collaboration, identifying potential issues, and ensuring adherence to coding standards for superior code quality.

Tata Consultancy Services, Chennai, India

November 2011 – May 2015

Senior Associate – Projects

- Coded and designed application solutions with Spring Boot and React.
- Implemented continuous data movement from archive tables to S3 buckets in AWS.
- Configured a usage-service endpoint in AWS API Gateway to consume data from a usage-capture system and provide responses to clients.
- Involved in performance tuning, query optimization, and code review at the end of the application development cycle.
- Utilized build-automation tools like Maven and implemented unit and integration tests using JUnit and TestNG to ensure code quality and reliability.